

Ensemble techniques -Project3 - Edouard Toutounji - 21 Feb 2020

October 27, 2025

1 1- Exploratory data analysis of the ‘bank-full.csv’ data as is - prior to any changes

1.1 1.1. Non visual analysis

Comments will follow to set a path on how to make necessary changes that will be implemented in part 2

```
[254]: # importing EDA libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
[255]: # loading the data into the dataframe TDS ( Term Deposit Sale)
TDS = pd.read_csv('bank-full.csv')
```

```
[256]: TDS.head(100)

# first impression : many of the 16 attributes are categorical
# and will require transformation to ordinal level or using dummy variables.
```

```
[256]:      age      job marital education default balance housing loan \
0      58 management married   tertiary    no     2143    yes    no
1      44 technician single   secondary   no      29    yes    no
2      33 entrepreneur married   secondary   no       2    yes    yes
3      47 blue-collar married   unknown    no     1506    yes    no
4      33           unknown single   unknown    no        1    no    no
..    ...
95     36 management married   tertiary    no      101    yes    yes
96     55 blue-collar married   secondary   no     383     no    no
97     60 retired   married   tertiary    no      81    yes    no
98     39 technician married   secondary   no       0    yes    no
99     46 management married   tertiary    no     229    yes    no

      contact day month duration campaign pdays previous poutcome Target
0    unknown    5   may       261         1      -1         0  unknown    no
```

```

1    unknown      5   may       151      1     -1      0  unknown    no
2    unknown      5   may        76      1     -1      0  unknown    no
3    unknown      5   may       92      1     -1      0  unknown    no
4    unknown      5   may      198      1     -1      0  unknown    no
...
95   unknown      5   may      426      1     -1      0  unknown    no
96   unknown      5   may      287      1     -1      0  unknown    no
97   unknown      5   may      101      1     -1      0  unknown    no
98   unknown      5   may      203      1     -1      0  unknown    no
99   unknown      5   may      197      1     -1      0  unknown    no

```

[100 rows x 17 columns]

[257]: # column names
TDS.columns

[257]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
 'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays',
 'previous', 'poutcome', 'Target'],
 dtype='object')

[258]: # data size and shape
TDS.shape

[258]: (45211, 17)

[259]: # dtypes of the attributes columns
TDS.dtypes

confirmation that many attributes are of type: 'object' --> to be transformed
 ↴ later

[259]: age int64
 job object
 marital object
 education object
 default object
 balance int64
 housing object
 loan object
 contact object
 day int64
 month object
 duration int64
 campaign int64
 pdays int64
 previous int64

```
poutcome      object  
Target        object  
dtype: object
```

```
[260]: TDS.info()  
# NO missing values however some surprises could be expected that may require  
↳replacement(s) later on.
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 45211 entries, 0 to 45210  
Data columns (total 17 columns):  
age            45211 non-null int64  
job            45211 non-null object  
marital         45211 non-null object  
education       45211 non-null object  
default          45211 non-null object  
balance          45211 non-null int64  
housing          45211 non-null object  
loan             45211 non-null object  
contact          45211 non-null object  
day              45211 non-null int64  
month            45211 non-null object  
duration         45211 non-null int64  
campaign         45211 non-null int64  
pdays            45211 non-null int64  
previous          45211 non-null int64  
poutcome         45211 non-null object  
Target            45211 non-null object  
dtypes: int64(7), object(10)  
memory usage: 5.9+ MB
```

```
[261]: # statistical analysis of the numerical attributes  
TDS.describe().transpose()  
  
# attributes where [ mean is quite > than median ] is an indicator of strong  
↳right skewness;  
# and presence of outliers in the respective distributions is also confirmed by  
↳the (max/min) values.  
  
# this is observed for : 'balance' , 'duration' , 'campaign' , 'pdays' ,  
↳'previous'
```

```
[261]:
```

	count	mean	std	min	25%	50%	75%	\
age	45211.0	40.936210	10.618762	18.0	33.0	39.0	48.0	
balance	45211.0	1362.272058	3044.765829	-8019.0	72.0	448.0	1428.0	
day	45211.0	15.806419	8.322476	1.0	8.0	16.0	21.0	
duration	45211.0	258.163080	257.527812	0.0	103.0	180.0	319.0	

```

campaign    45211.0      2.763841      3.098021      1.0      1.0      2.0      3.0
pdays       45211.0     40.197828    100.128746     -1.0     -1.0     -1.0     -1.0
previous    45211.0      0.580323      2.303441      0.0      0.0      0.0      0.0

                           max
age             95.0
balance        102127.0
day              31.0
duration        4918.0
campaign        63.0
pdays           871.0
previous        275.0

```

[262]: # confirming our hunch about the data by printing the values counts of all columns/attributes.

```

for column in TDS.columns:
    print(TDS[column].value_counts(), '\n')

```

```

32      2085
31      1996
33      1972
34      1930
35      1894
...
90      2
92      2
93      2
95      2
94      1
Name: age, Length: 77, dtype: int64

```

```

blue-collar      9732
management       9458
technician        7597
admin.          5171
services         4154
retired          2264
self-employed    1579
entrepreneur     1487
unemployed       1303
housemaid        1240
student          938
unknown          288
Name: job, dtype: int64

```

```

married        27214
single         12790

```

```
divorced      5207  
Name: marital, dtype: int64
```

```
secondary     23202  
tertiary      13301  
primary       6851  
unknown       1857  
Name: education, dtype: int64
```

```
no          44396  
yes         815  
Name: default, dtype: int64
```

```
0            3514  
1            195  
2            156  
4            139  
3            134  
...  
4305         1  
6352         1  
18881        1  
14889        1  
7218         1  
Name: balance, Length: 7168, dtype: int64
```

```
yes        25130  
no         20081  
Name: housing, dtype: int64
```

```
no        37967  
yes       7244  
Name: loan, dtype: int64
```

```
cellular    29285  
unknown     13020  
telephone   2906  
Name: contact, dtype: int64
```

```
20        2752  
18        2308  
21        2026  
17        1939  
6         1932  
5         1910  
14        1848  
8         1842  
28        1830
```

```
7      1817
19     1757
29     1745
15     1703
12     1603
13     1585
30     1566
9      1561
11     1479
4      1445
16     1415
2      1293
27     1121
3      1079
26     1035
23     939
22     905
25     840
31     643
10     524
24     447
1      322
Name: day, dtype: int64
```

```
may    13766
jul     6895
aug     6247
jun     5341
nov     3970
apr     2932
feb     2649
jan     1403
oct     738
sep     579
mar     477
dec     214
Name: month, dtype: int64
```

```
124     188
90      184
89      177
122     175
104     175
...
2150     1
1970     1
1906     1
1842     1
```

```
2015      1  
Name: duration, Length: 1573, dtype: int64
```

```
1    17544  
2    12505  
3    5521  
4    3522  
5    1764  
6    1291  
7    735  
8    540  
9    327  
10   266  
11   201  
12   155  
13   133  
14   93  
15   84  
16   79  
17   69  
18   51  
19   44  
20   43  
21   35  
22   23  
23   22  
25   22  
24   20  
28   16  
29   16  
26   13  
31   12  
27   10  
32   9  
30   8  
33   6  
34   5  
36   4  
35   4  
43   3  
38   3  
41   2  
50   2  
37   2  
51   1  
55   1  
46   1  
58   1
```

```
44      1  
39      1  
63      1  
Name: campaign, dtype: int64
```

```
-1      36954  
182     167  
92      147  
183     126  
91      126  
...  
749     1  
717     1  
589     1  
493     1  
32      1
```

```
Name: pdays, Length: 559, dtype: int64
```

```
0      36954  
1      2772  
2      2106  
3      1142  
4      714  
5      459  
6      277  
7      205  
8      129  
9      92  
10     67  
11     65  
12     44  
13     38  
15     20  
14     19  
17     15  
16     13  
19     11  
23     8  
20     8  
22     6  
18     6  
24     5  
27     5  
29     4  
25     4  
21     4  
30     3  
28     2
```

```
26      2
37      2
38      2
55      1
40      1
35      1
58      1
51      1
41      1
32      1
275     1
Name: previous, dtype: int64
```

```
unknown    36959
failure     4901
other       1840
success     1511
Name: poutcome, dtype: int64
```

```
no      39922
yes     5289
Name: Target, dtype: int64
```

Before jumping to graphical univariate analysis , some important preliminary observations on the distributions of some attributes.

- 1 - ‘Target’ : is tilted largely towards ‘no’. (so the chances of guessing by chance are not 50 - 50)
- 2 - ‘poutcome’: shows that clients labeled as ‘unknown’ are = 36959
- 3 - ‘pdays’ : shows that clients labeled as ‘-1’ = 36954

This seems to confirm that previously uncontacted clients amount to approx. 36950.

- 4 - ‘previous’: shows strong right skewness and an outlier ‘275’
- 5 - ‘campaign’: shows strong right skewness and an outlier ‘63’

1.2 1.2. Univariate Graphical analysis column by column.


```
[263]: TDS.columns
```

```
[263]: Index(['age', 'job', 'marital', 'education', 'default', 'balance', 'housing',
       'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'Target'],
       dtype='object')
```

```
[264]: # plotting the distributions of All 17 columns.
```


```
sns.boxplot(TDS.age)
```

```
[264]: <matplotlib.axes._subplots.AxesSubplot at 0x1a43724090>
```



```
[265]: sns.distplot(TDS.age)
# boxplot for 'age' - a big drop after 60.
```

```
[265]: <matplotlib.axes._subplots.AxesSubplot at 0x1a44511390>
```




```
[266]: sns.countplot(TDS.job)
plt.xticks(rotation=45, horizontalalignment='right')

# 'unknown' category could be dropped as it gives no insight into classifying
# clients and is minimal in count.

TDS.job.value_counts()
```


```
[266]: blue-collar      9732
       management       9458
       technician        7597
       admin.            5171
       services          4154
       retired           2264
       self-employed     1579
       entrepreneur      1487
       unemployed        1303
       housemaid         1240
       student            938
       unknown             288
Name: job, dtype: int64
```



```
[267]: sns.countplot(TDS.marital)
plt.xticks(rotation=45, horizontalalignment='right')

# majority of our pool are 'married'
TDS.marital.value_counts()
```


```
[267]: married      27214
       single       12790
       divorced     5207
Name: marital, dtype: int64
```



```
[268]: sns.countplot(TDS.education)
plt.xticks(rotation=45, horizontalalignment='right')

# 'unknown' education label is a flag for later analysis
TDS.education.value_counts()
```


```
[268]: secondary      23202
tertiary        13301
primary         6851
unknown         1857
Name: education, dtype: int64
```



```
[269]: sns.countplot(TDS.default)
plt.xticks(rotation=45, horizontalalignment='right')

# imbalance : the Big majority have not defaulted. There are no unknowns.
TDS.default.value_counts()
```


```
[269]: no      44396
yes      815
Name: default, dtype: int64
```



```
[270]: sns.distplot(TDS.balance)
plt.xticks(rotation=45, horizontalalignment='right')

# outliers could be those clients with balance above 20,000.
```


```
[270]: (array([-20000.,      0.,  20000.,  40000.,  60000.,  80000., 100000.,
       120000.]), <a list of 8 Text xticklabel objects>)
```



```
[271]: sns.countplot(TDS.housing)
plt.xticks(rotation=45, horizontalalignment='right')

# no strong influence expected from housing loans as it is roughly balanced in
# our pool
TDS.housing.value_counts()
```


```
[271]: yes    25130
no     20081
Name: housing, dtype: int64
```



```
[272]: sns.countplot(TDS.loan)
plt.xticks(rotation=45, horizontalalignment='right')

# Majority do not have personal loans
TDS.loan.value_counts()
```


```
[272]: no      37967
yes     7244
Name: loan, dtype: int64
```



```
[273]: sns.countplot(TDS.contact)
plt.xticks(rotation=45, horizontalalignment='right')

# Majority have cell phones
# Unknowns will be flagged for later consideration
TDS.contact.value_counts()
```


```
[273]: cellular      29285
unknown        13020
telephone       2906
Name: contact, dtype: int64
```



```
[274]: sns.countplot(TDS.day)
plt.xticks(rotation=45, horizontalalignment='right')

# the day of the month is not equiprobable but spreadout. Pyschological factors at play ?
# few days of the month seem to be lower than others:
...
23      939
22      905
25      840
31      643
10      524
24      447
1       322
...
```


[274]: '\n23 939\n22 905\n25 840\n31 643\n10 524\n24 447\n1 322\n'



```
[275]: sns.countplot(TDS.month)
plt.xticks(rotation=45, horizontalalignment='right')

# Obviously some months are more active than others for contacts with clients.
# Maybe it is a limited promotional timing that was mainly launched in May.
TDS.month.value_counts()
```


```
[275]: may      13766
       jul      6895
       aug      6247
       jun      5341
       nov      3970
       apr      2932
       feb      2649
       jan      1403
       oct      738
       sep      579
       mar      477
       dec      214
Name: month, dtype: int64
```



```
[276]: sns.distplot(TDS.duration)
plt.xticks(rotation=45, horizontalalignment='right')

# A right skewed distribution with many outliers, most calls are below 1000
# seconds in length.
# but it was mentionned that it's preferable to drop it out of the model later
# on.
# It could be that the longest calls resulted in positive outcomes for TDS so
# we cannot make a cutoff .
```


```
[276]: (array([-1000., 0., 1000., 2000., 3000., 4000., 5000., 6000.]),
<a list of 8 Text xticklabel objects>)
```



```
[277]: sns.distplot(TDS.campaign)
plt.xticks(rotation=45, horizontalalignment='right')

# also 'campaign' very right skewed and could be correlated with 'duration'
```


```
[277]: (array([-10.,  0.,  10.,  20.,  30.,  40.,  50.,  60.,  70.]),
 <a list of 9 Text xticklabel objects>)
```



```
[278]: sns.countplot(TDS.previous)
plt.xticks(rotation=45, horizontalalignment='right')

# also 'previous' campaign very right skewed and could be correlated with
# 'duration'
```


```
[278]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
       34, 35, 36, 37, 38, 39, 40]), <a list of 41 Text xticklabel objects>)
```



```
[279]: sns.distplot(TDS.pdays)
plt.xticks(rotation=45, horizontalalignment='right')

# '-1' is the majority who were not contacted
TDS.pdays.value_counts()
```


```
[279]: -1      36954  
        182     167  
        92      147  
       183     126  
       91      126  
       ...  
       749      1  
       717      1  
       589      1  
       493      1  
       32       1  
Name: pdays, Length: 559, dtype: int64
```



```
[280]: sns.countplot(TDS.poutcome)
plt.xticks(rotation=45, horizontalalignment='right')

# 'unknown' is the majority who were not contacted in the previous campaign. ↴
# correlated with the '-1' value above.
TDS.poutcome.value_counts()
```


```
[280]: unknown      36959
failure        4901
other          1840
success        1511
Name: poutcome, dtype: int64
```



```
[281]: sns.countplot(TDS.Target)
plt.xticks(rotation=45, horizontalalignment='right')

# vast majority did not subscribe. Is it due to failure of being contact in a
# previous campaign ?
TDS.Target.value_counts()
```

```
[281]: no      39922
yes     5289
Name: Target, dtype: int64
```



2 2- Preparing Data for Analytics

Transforming attributes that are object or categorical and then proceed with joinplots to explore correlations

```
[282]: TDS.dtypes
```

```
[282]: age           int64
job            object
marital        object
education      object
default         object
balance        int64
housing         object
loan            object
contact         object
day             int64
month           object
duration       int64
campaign        int64
pdays           int64
previous        int64
poutcome        object
```

```
Target          object
dtype: object
```

```
[283]: # transforming object types columns into categorical
```

```
for column in TDS.columns:
    if TDS[column].dtype == 'object':
        TDS[column] = pd.Categorical(TDS[column])
```

```
[284]: TDS.dtypes
# :-)
```

```
[284]: age           int64
job            category
marital        category
education      category
default         category
balance         int64
housing         category
loan            category
contact         category
day             int64
month           category
duration        int64
campaign        int64
pdays           int64
previous        int64
poutcome        category
Target          category
dtype: object
```

```
[327]: # now we have to use One-Hot encoding for the columns: 'job' , 'marital', ↴ 'month'.
```

```
TDS_dummies = pd.get_dummies(TDS, columns= ['job' , 'marital', 'month'])

TDS_dummies.head()
```

```
[327]:   age  education default  balance housing loan  contact  day  duration \
0    58    tertiary     no     2143    yes    no  unknown    5    261
1    44  secondary     no      29    yes    no  unknown    5    151
2    33  secondary     no       2    yes   yes  unknown    5     76
3    47    unknown     no     1506    yes    no  unknown    5     92
4    33    unknown     no       1    no    no  unknown    5    198

  campaign ... month_dec  month_feb month_jan month_jul month_jun \
0          1   ...        0          0          0          0
```

```

1      1 ...      0      0      0      0
2      1 ...      0      0      0      0
3      1 ...      0      0      0      0
4      1 ...      0      0      0      0

    month_mar  month_may  month_nov  month_oct  month_sep
0          0         1         0         0         0
1          0         1         0         0         0
2          0         1         0         0         0
3          0         1         0         0         0
4          0         1         0         0         0

```

[5 rows x 41 columns]

[348]: TDS_dummies.Target.value_counts()

[348]: no 39922
yes 5289
Name: Target, dtype: int64

[349]: # Changing remaining categorical columns into ordinal scale & Creating new DF :
 ↪TDS_dummies_repl

```

replaceStruct = {
    'education': {'unknown': 0, 'primary': 1, 'secondary': 2,
    ↪, 'tertiary':3},
    'default' : {'yes': 0, 'no':1 },
    # defaulting is not a good indicator hence 'yes'
    ↪--> 0
    'housing' : {'no':1 , 'yes':0},
    'loan' : {'no':1 , 'yes':0},
    'contact' : {'unknown':0, 'telephone':1, 'cellular':2},
    'poutcome' : {'failure':-1,'unknown':0,'other':0, 'success':1},
    # unknown and other considered same neutral level
    'Target' : {'no':0 , 'yes':1},
    # unknown and other considered same neutral level
    'pdays' : {-1: 999}
    #replacing -1 by 999 to show that client was never
    ↪contacted
}
TDS_dummies_repl = TDS_dummies.replace(replaceStruct)

TDS_dummies_repl.head()

# All columns are numerical now :-) and hence correlations are possible!

```

```
[349]:    age education default balance housing loan contact day duration \
0      58          3       1     2143      0      1      0      5     261
1      44          2       1      29      0      1      0      5     151
2      33          2       1       2      0      0      0      5      76
3      47          0       1    1506      0      1      0      5      92
4      33          0       1       1      1      1      0      5    198

    campaign ... month_dec month_feb month_jan month_jul month_jun \
0           1   ...        0        0        0        0        0
1           1   ...        0        0        0        0        0
2           1   ...        0        0        0        0        0
3           1   ...        0        0        0        0        0
4           1   ...        0        0        0        0        0

    month_mar month_may month_nov month_oct month_sep
0         0        1        0        0        0
1         0        1        0        0        0
2         0        1        0        0        0
3         0        1        0        0        0
4         0        1        0        0        0

[5 rows x 41 columns]
```

```
[331]: TDS_dummies_repl.Target.value_counts()
```

```
[331]: 0    39922
1    5289
Name: Target, dtype: int64
```

```
[332]: # SCALING numerical columns: age, balance, duration, campaign, pdays, previous

for col in TDS_dummies_repl.columns:

    if col == 'age':
        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())
    elif col == 'balance':
        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())
    elif col == 'duration':
        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())
    elif col == 'campaign':
        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())
    elif col == 'pdays':
```

```

        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())
    elif col == 'previous':
        TDS_dummies_repl[col] = TDS_dummies_repl[col]/(TDS_dummies_repl[col].
        ↪max())

TDS_dummies_repl.head()

```

```
[332]:      age education default balance housing loan contact day \
0 0.610526            3       1 0.020984      0   1     0   5
1 0.463158            2       1 0.000284      0   1     0   5
2 0.347368            2       1 0.000020      0   0     0   5
3 0.494737            0       1 0.014746      0   1     0   5
4 0.347368            0       1 0.000010      1   1     0   5

duration campaign ... month_dec month_feb month_jan month_jul \
0 0.053070 0.015873 ...          0          0          0          0
1 0.030704 0.015873 ...          0          0          0          0
2 0.015453 0.015873 ...          0          0          0          0
3 0.018707 0.015873 ...          0          0          0          0
4 0.040260 0.015873 ...          0          0          0          0

month_jun month_mar month_may month_nov month_oct month_sep
0          0          0          1          0          0          0
1          0          0          1          0          0          0
2          0          0          1          0          0          0
3          0          0          1          0          0          0
4          0          0          1          0          0          0
```

[5 rows x 41 columns]


```
[311]: plt.figure(figsize=(20,20))

sns.heatmap( TDS_dummies_repl.corr(),
              annot=True,
              linewidths=.5,
              center=0,
              cbar=False,
              cmap="YlGnBu")

# NO Obvious Proxies with a correlation > 0.9 !

```

[311]: <matplotlib.axes._subplots.AxesSubplot at 0x1a4a133690>



```
[1]: # NOTE : tried sns.pairplot( TDS_dummies_repl ) ... it was too computationally heavy.
#couldn't get graphical results
```

3 3 and 4 - DT and Ensemble Models and comparison of performance

```
[368]: # Importing the needed sklearn packages
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, recall_score,
precision_score, confusion_matrix
```

```
from sklearn.tree      import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
```

[350]: TDS_dummies_repl.info()

```
# all columns are numerical types
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 41 columns):
age            45211 non-null int64
education      45211 non-null int64
default         45211 non-null int64
balance         45211 non-null int64
housing          45211 non-null int64
loan             45211 non-null int64
contact          45211 non-null int64
day              45211 non-null int64
duration         45211 non-null int64
campaign         45211 non-null int64
pdays            45211 non-null int64
previous          45211 non-null int64
poutcome          45211 non-null int64
Target            45211 non-null int64
job_admin.        45211 non-null uint8
job_blue-collar   45211 non-null uint8
job_entrepreneur    45211 non-null uint8
job_housemaid       45211 non-null uint8
job_management        45211 non-null uint8
job_retired          45211 non-null uint8
job_self-employed     45211 non-null uint8
job_services          45211 non-null uint8
job_student           45211 non-null uint8
job_technician         45211 non-null uint8
job_unemployed         45211 non-null uint8
job_unknown            45211 non-null uint8
marital_divorced       45211 non-null uint8
marital_married        45211 non-null uint8
marital_single          45211 non-null uint8
month_apr              45211 non-null uint8
month_aug              45211 non-null uint8
month_dec              45211 non-null uint8
month_feb              45211 non-null uint8
month_jan              45211 non-null uint8
month_jul              45211 non-null uint8
```

```

month_jun          45211 non-null uint8
month_mar         45211 non-null uint8
month_may         45211 non-null uint8
month_nov         45211 non-null uint8
month_oct         45211 non-null uint8
month_sep         45211 non-null uint8
dtypes: int64(14), uint8(27)
memory usage: 6.0 MB

```

3.1 3.1 split data

```
[399]: X = TDS_dummies_repl.drop( ["Target", "duration"], axis=1)
y = TDS_dummies_repl.Target
```

```
[400]: print(X.shape)
print(y.shape)
```

```
(45211, 39)
(45211,)
```

```
[441]: (X_train, X_test, y_train, y_test) = train_test_split(X, y, test_size=.30, random_state=1)

print(X_test.shape)
print(y_test.shape)
```

```
(13564, 39)
(13564,)
```

3.2 3.2 Decision Tree model using GINI criteria

```
[402]: dTree_Gini= DecisionTreeClassifier(criterion = 'gini', random_state=1)

dTree_Gini.fit(X_train, y_train)

print(dTree_Gini.score(X_train, y_train))
print(dTree_Gini.score(X_test, y_test))

# The expected overfitting is there with the training Data
```

```
1.0
0.8352255971689767
```

```
[461]: y_pred_G = dTree_Gini.predict(X_test)

# CONFUSION matrix
cm = metrics.confusion_matrix(y_test, y_pred_G, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No", "Yes"]],
```

```

        columns = [i for i in ["No", "Yes"]])

df_cm
```

[461]:

	No	Yes
No	10787	1226
Yes	1009	542

[462]:

```

print(' Accuracy - Recall - Precision scores respectively: ')
print(metrics.accuracy_score(y_test, y_pred_G))
print(metrics.recall_score(y_test, y_pred_G))
print(metrics.precision_score(y_test, y_pred_G))
```

Accuracy - Recall - Precision scores respectively:
0.8352255971689767
0.34945196647324306
0.3065610859728507

[404]:

```

print(pd.DataFrame(dTree_Gini.feature_importances_, columns = ["Imp"],
                   index=X_train.columns))

# 'balance', 'age', 'day' seems to be the most important features
```

	Imp
age	0.143323
education	0.031835
default	0.002503
balance	0.204176
housing	0.024161
loan	0.012849
contact	0.013699
day	0.120806
campaign	0.054560
pdays	0.046922
previous	0.024002
poutcome	0.098443
job_admin.	0.010719
job_blue-collar	0.012416
job_entrepreneur	0.004377
job_housemaid	0.004178
job_management	0.013014
job_retired	0.005220
job_self-employed	0.006320
job_services	0.010259
job_student	0.005797
job_technician	0.017345
job_unemployed	0.006367
job_unknown	0.000862

```

marital_divorced    0.005912
marital_married     0.009750
marital_single      0.007906
month_apr           0.007355
month_aug           0.009540
month_dec           0.004198
month_feb           0.009696
month_jan           0.004449
month_jul           0.007358
month_jun           0.012728
month_mar           0.011076
month_may           0.007761
month_nov           0.006384
month_oct           0.011306
month_sep           0.010429

```

3.3 Decision Tree model using GINI criteria and depth = 4 (Pruning)

```
[463]: dTree_Gini_4= DecisionTreeClassifier(criterion = 'gini', max_depth=4
                                         ,random_state=1)
```

```

dTree_Gini_4.fit(X_train, y_train)

print(dTree_Gini_4.score(X_train, y_train))
print(dTree_Gini_4.score(X_test, y_test))

# Testing results are much better than without pruning
```

0.8942395803709672

0.8951636685343557

```
[464]: y_pred_G4 = dTree_Gini_4.predict(X_test)
```

```

# CONFUSION matrix

cm = metrics.confusion_matrix(y_test, y_pred_G4, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No", "Yes"]],
                     columns = [i for i in ["No", "Yes"]])

df_cm
```

	No	Yes
No	11828	185
Yes	1237	314

```
[465]: print(' Accuracy - Recall - Precision scores respectively: ')
print(metrics.accuracy_score(y_test, y_pred_G4))
```

```
print(metrics.recall_score(y_test, y_pred_G4))
print(metrics.precision_score(y_test, y_pred_G4))
```

Accuracy - Recall - Precision scores respectively:

0.8951636685343557

0.2024500322372663

0.6292585170340681

```
[419]: print(pd.DataFrame(dTree_Gini_4.feature_importances_, columns = ["Imp"],
                           index=X_train.columns))
```

'poutcome' seems to be the most important feature by large!

	Imp
age	0.110006
education	0.000000
default	0.000000
balance	0.000000
housing	0.016810
loan	0.000000
contact	0.000000
day	0.024734
campaign	0.002670
pdays	0.016754
previous	0.000000
poutcome	0.661688
job_admin.	0.000000
job_blue-collar	0.000000
job_entrepreneur	0.000000
job_housemaid	0.000000
job_management	0.000000
job_retired	0.000000
job_self-employed	0.000000
job_services	0.000000
job_student	0.000000
job_technician	0.002572
job_unemployed	0.000000
job_unknown	0.000000
marital_divorced	0.004168
marital_married	0.000000
marital_single	0.000000
month_apr	0.000000
month_aug	0.000000
month_dec	0.004339
month_feb	0.000000
month_jan	0.000000
month_jul	0.000000
month_jun	0.000000

```
month_mar      0.074116
month_may      0.007426
month_nov      0.000000
month_oct      0.074716
month_sep      0.000000
```

3.4 Decision Tree model using Entropy criteria

```
[475]: dTree_Entropy= DecisionTreeClassifier(criterion = 'entropy' ,random_state=1)

dTree_Entropy.fit(X_train, y_train)

print(dTree_Entropy.score(X_train, y_train))
print(dTree_Entropy.score(X_test, y_test))

# Overfitting expected with training data
```

```
1.0
0.8347095252138013
```

```
[476]: y_pred_E= dTree_Entropy.predict(X_test)

# CONFUSION matrix

cm = metrics.confusion_matrix(y_test, y_pred_E, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No","Yes"]],  
                     columns = [i for i in ["No","Yes"]])

df_cm
```

```
[476]:      No    Yes
No     10828   1185
Yes     1057    494
```

```
[477]: print(' Accuracy - Recall - Precision scores respectively: ')
print(metrics.accuracy_score(y_test, y_pred_E))
print(metrics.recall_score(y_test, y_pred_E))
print(metrics.precision_score(y_test, y_pred_E))
```

```
Accuracy - Recall - Precision scores respectively:
0.8347095252138013
0.3185041908446164
0.2942227516378797
```

```
[422]: print (pd.DataFrame(dTree_Entropy.feature_importances_, columns = ["Imp"],  
                           index = X_train.columns))
```

```
# 'balance', 'age' , 'day' seem to be the most important features
```

	Imp
age	0.144494
education	0.036393
default	0.002601
balance	0.219630
housing	0.016815
loan	0.013681
contact	0.030039
day	0.112352
campaign	0.059784
pdays	0.046608
previous	0.019791
poutcome	0.081725
job_admin.	0.013046
job_blue-collar	0.015027
job_entrepreneur	0.004365
job_housemaid	0.002829
job_management	0.013128
job_retired	0.003448
job_self-employed	0.006259
job_services	0.010182
job_student	0.003707
job_technician	0.011517
job_unemployed	0.005701
job_unknown	0.002033
marital_divorced	0.007445
marital_married	0.009797
marital_single	0.008096
month_apr	0.012160
month_aug	0.007104
month_dec	0.003841
month_feb	0.008604
month_jan	0.004092
month_jul	0.006509
month_jun	0.011545
month_mar	0.009458
month_may	0.008560
month_nov	0.007789
month_oct	0.010675
month_sep	0.009170

3.5 Decision Tree model using Entropy criteria and depth = 4 (Pruning)

```
[423]: dTree_Entropy_4= DecisionTreeClassifier(criterion = 'entropy', max_depth=4,random_state=1)
```

```
dTree_Entropy_4.fit(X_train, y_train)
```

```
print(dTree_Entropy_4.score(X_train, y_train))  
print(dTree_Entropy_4.score(X_test, y_test))
```

```
# Testing results are much better than without pruning
```

```
0.8927228489272285
```

```
0.8956797404895311
```

```
[478]: y_pred_E4 = dTree_Entropy_4.predict(X_test)
```

```
# CONFUSION matrix
```

```
cm = metrics.confusion_matrix(y_test, y_pred_E4, labels=[0, 1])
```

```
df_cm = pd.DataFrame(cm, index = [i for i in ["No","Yes"]],  
columns = [i for i in ["No","Yes"]])
```

```
df_cm
```

```
[478]:
```

	No	Yes
No	11866	147
Yes	1268	283

```
[479]: print(' Accuracy - Recall - Precision scores respectively: ')
```

```
print(metrics.accuracy_score(y_test, y_pred_E4))
```

```
print(metrics.recall_score(y_test, y_pred_E4))
```

```
print(metrics.precision_score(y_test, y_pred_E4))
```

```
Accuracy - Recall - Precision scores respectively:
```

```
0.8956797404895311
```

```
0.1824629271437782
```

```
0.6581395348837209
```

```
[425]: print(pd.DataFrame(dTree_Entropy_4.feature_importances_, columns = ["Imp"],  
index = X_train.columns))
```

```
# 'poutcome' seems to be the most important feature by far
```

	Imp
age	0.096045
education	0.000000
default	0.000000

balance	0.000000
housing	0.011200
loan	0.000000
contact	0.184645
day	0.014742
campaign	0.000000
pdays	0.009188
previous	0.000000
poutcome	0.584279
job_admin.	0.000000
job_blue-collar	0.000000
job_entrepreneur	0.000000
job_housemaid	0.000000
job_management	0.000000
job_retired	0.000000
job_self-employed	0.000000
job_services	0.000000
job_student	0.000000
job_technician	0.002000
job_unemployed	0.000000
job_unknown	0.000000
marital_divorced	0.000000
marital_married	0.008561
marital_single	0.000000
month_apr	0.000000
month_aug	0.000000
month_dec	0.000000
month_feb	0.000000
month_jan	0.000000
month_jul	0.000000
month_jun	0.074695
month_mar	0.000000
month_may	0.004646
month_nov	0.000000
month_oct	0.010000
month_sep	0.000000

4 Ensemble Methods

4.1 3.6 Random Forest model

```
[471]: RFcl = RandomForestClassifier(n_estimators = 10, random_state=1,max_features=12)
RFcl = RFcl.fit(X_train, y_train)

print(RFcl.score(X_train, y_train))
print(RFcl.score(X_test, y_test))
```

0.9855594527127374

```
0.8913299911530522
```

```
[472]: y_pred_RF = RFcl.predict(X_test)

# CONFUSION matrix

cm = metrics.confusion_matrix(y_test, y_pred_RF, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No", "Yes"]],  
                     columns = [i for i in ["No", "Yes"]])

df_cm
```

```
[472]:      No   Yes  
No    11746   267  
Yes     1207   344
```

```
[473]: print(' Accuracy - Recall - Precision scores respectively: ')  
print(metrics.accuracy_score(y_test, y_pred_RF))  
print(metrics.recall_score(y_test, y_pred_RF))  
print(metrics.precision_score(y_test, y_pred_RF))
```

```
Accuracy - Recall - Precision scores respectively:  
0.8913299911530522  
0.22179239200515796  
0.563011456628478
```

4.2 3.7 Adaboost Ensemble model

```
[468]: ADcl = AdaBoostClassifier(n_estimators=10, random_state=1)

ADcl = ADcl.fit(X_train, y_train)

print(ADcl.score(X_train, y_train))
print(ADcl.score(X_test, y_test))
```

```
0.8912061174834898  
0.8942789737540549
```

```
[469]: y_pred_AD = ADcl.predict(X_test)

# CONFUSION matrix

cm = metrics.confusion_matrix(y_test, y_pred_AD, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No", "Yes"]],  
                     columns = [i for i in ["No", "Yes"]])

df_cm
```

```
[469]:      No  Yes
No    11809  204
Yes    1230  321
```

```
[470]: print(' Accuracy - Recall - Precision scores respectively: ')
print(metrics.accuracy_score(y_test, y_pred_AD))
print(metrics.recall_score(y_test, y_pred_AD))
print(metrics.precision_score(y_test, y_pred_AD))
```

```
Accuracy - Recall - Precision scores respectively:
0.8942789737540549
0.20696324951644102
0.6114285714285714
```

4.3 3.7 Bagging Ensemble model

```
[449]: BGcl = BaggingClassifier(n_estimators=10,random_state=1)

BGcl = BGcl.fit(X_train, y_train)

print(BGcl.score(X_train, y_train))
print(BGcl.score(X_test, y_test))
```

```
0.9857174455714601
0.8900766735476261
```

```
[466]: y_pred_BG = bgcl.predict(X_test)

# CONFUSION matrix

cm = metrics.confusion_matrix(y_test, y_pred_BG, labels=[0, 1])

df_cm = pd.DataFrame(cm, index = [i for i in ["No","Yes"]],
                      columns = [i for i in ["No","Yes"]])

df_cm
```

```
[466]:      No  Yes
No    11696  317
Yes    1174  377
```

```
[467]: print(' Accuracy - Recall - Precision scores respectively: ')
print(metrics.accuracy_score(y_test, y_pred_BG))
print(metrics.recall_score(y_test, y_pred_BG))
print(metrics.precision_score(y_test, y_pred_BG))
```

```
Accuracy - Recall - Precision scores respectively:
0.8900766735476261
```

0.24306898774983882

0.5432276657060519

5 4-Conclusion :

- 5.0.1 Most models - whether a simple pruned DT or Ensemble methods - give similar performance on the Testing Data with Accuracy Scores approx = 90%.
- 5.0.2 However Precision and Recall scores are not so good for most models although there are differences.
- 5.0.3 For the Ensemble methods , I did not provide a base model and used the default one , since even poor learners when aggregated give good predictions. I kept the number of Trees to 10 for comparison purposes and reduction of computational load as more did not provide significant improvements in prediction capabilities .

Project 3 - Ensemble Techniques & Model Tuning

Thank you - Edouard Toutounji - 21 Feb 2020 .

[]: